

2019  
怪兽  
学堂



# Gradient descent and optimizer

虾米 2019-3

# Gradient Descent for Linear Regression

Goal: minimize the following  
loss function:

predict with:  $\hat{y}^i = \sum_j^n w_j \phi_j(\mathbf{x}^i)$

$$J_{\mathbf{x},\mathbf{y}}(\mathbf{w}) = \sum_i (\mathcal{Y}^i - \hat{\mathcal{Y}}^i)^2 = \sum_i \left( \mathcal{Y}^i - \sum_j w_j \phi_j(\mathbf{x}^i) \right)^2$$

↑  
sum over  $n$  examples

↑  
sum over  $k+1$  basis vectors

# Gradient Descent for Linear Regression

Goal: minimize the following loss function:

predict with:  $\hat{y}^i = \sum_j^n w_j \phi_j(\mathbf{x}^i)$

$$J_{\mathbf{x},\mathbf{y}}(\mathbf{w}) = \sum_i (\mathcal{Y}^i - \hat{\mathcal{Y}}^i)^2 = \sum_i \left( \mathcal{Y}^i - \sum_j w_j \phi_j(\mathbf{x}^i) \right)^2$$

$$\frac{\partial}{\partial w_j} J(\mathbf{w}) = \frac{\partial}{\partial w_j} \sum_i (\mathcal{Y}^i - \hat{\mathcal{Y}}^i)^2$$
$$= 2 \sum_i (\mathcal{Y}^i - \hat{\mathcal{Y}}^i) \frac{\partial}{\partial w_j} \hat{\mathcal{Y}}^i$$

$$= 2 \sum_i (\mathcal{Y}^i - \hat{\mathcal{Y}}^i) \frac{\partial}{\partial w_j} \sum_j w_j \phi_j(\mathbf{x}^i)$$

$$= 2 \sum_i (\mathcal{Y}^i - \hat{\mathcal{Y}}^i) \phi_j(\mathbf{x}^i)$$





# Gradient Descent for Linear Regression

Learning algorithm:

- Initialize weights **w=0**
- For t=1,... until convergence:
  - Predict for each example  $\mathbf{x}^i$  using **w**:  $\hat{y}^i = \sum_{j=0}^k w_j \phi_j(\mathbf{x}^i)$
  - Compute gradient of loss:  $\frac{\partial}{\partial w_j} J(\mathbf{w}) = 2 \sum_i (y^i - \hat{y}^i) \phi_j(\mathbf{x}^i)$
  - This is a vector **g**
  - Update: **w = w - λg**
  - λ is the learning rate.

# Linear regression is a *convex* optimization problem





so again gradient descent will reach a *global* optimum



proof: differentiate again to get the second derivative

# Some issues about Gradient Descent

Learning algorithm:

- Initialize weights  **$\mathbf{w}=0$**  
- For  $t=1, \dots$  until convergence:
  - Predict for each example  $\mathbf{x}^i$  using  $\mathbf{w}$ :  $\hat{y}^i = \sum_{j=0}^k w_j \phi_j(\mathbf{x}^i)$
  - Compute gradient of loss:  $\frac{\partial}{\partial w_j} J(\mathbf{w}) = 2 \sum_i (y^i - \hat{y}^i) \phi_j(\mathbf{x}^i)$
  - This is a vector  **$\mathbf{g}$**  
  - Update:  **$\mathbf{w} = \mathbf{w} - \lambda \mathbf{g}$** ,  $\lambda$  is the learning rate.   


Initial status


Gradient computation

Weight update

Learning rate setting



# Weight initialization status




# Weight initialization status

- An incorrect initialization can slow down or even completely stall the learning process. Luckily, this issue can be diagnosed relatively easily
- One way to do so is to plot activation/gradient histograms for all layers of the network.




# Weight initialization status

- First-layer Visualizations



Noisy features indicate could be a symptom: Unconverged network, improperly set learning rate, very low weight regularization penalty



Nice, smooth, clean and diverse features are a good indication that the training is proceeding well.

# Gradient computation

- Gradient Checks
- In theory, performing a gradient check is as simple as comparing the analytic gradient to the numerical gradient.

$$\frac{df(x)}{dx} = \frac{f(x+h) - f(x)}{h} \quad (\text{bad, do not use})$$

$$\frac{df(x)}{dx} = \frac{f(x+h) - f(x-h)}{2h} \quad (\text{use instead})$$

$h$  is a very small number, in practice approximately  $1e-5$

In practice, it turns out that it is much better to use the *centered difference formula*

# Gradient computation:

## Stochastic Gradient Descent

$$L = \sum_n \left( y^n - \left( b + \sum w_i x_i^n \right) \right)^2$$

Loss is the summation over all training examples

- **Gradient Descent**  $\theta^i = \theta^{i-1} - \eta \nabla L(\theta^{i-1})$
- **Stochastic Gradient Descent**

Faster!

Pick an example  $x^n$

$$L^n = \left( y^n - \left( b + \sum w_i x_i^n \right) \right)^2 \quad \theta^i = \theta^{i-1} - \eta \nabla L^n(\theta^{i-1})$$

Loss for only one example




# Gradient computation:

## Stochastic Gradient Descent


### *Gradient Descent*

Update after seeing all examples



### *Stochastic Gradient Descent*

Update for each example



If there are 20 examples, 20 times faster.

# Gradient computation:

## Stochastic Gradient Descent

Gradient  
Descent

Stochastic  
Gradient Descent

Batch based  
Gradient Descent




```
graph TD; A[Gradient Descent] --> C[Batch based Gradient Descent]; B[Stochastic Gradient Descent] --> C;
```

The diagram illustrates the relationship between three types of gradient descent. At the top, 'Gradient computation:' is followed by 'Stochastic Gradient Descent' in blue. Below this, 'Gradient Descent' and 'Stochastic Gradient Descent' are positioned on the left and right respectively. Two orange arrows point from both 'Gradient Descent' and 'Stochastic Gradient Descent' towards 'Batch based Gradient Descent' at the bottom center, indicating that both are specific instances or variations of the general gradient descent concept.

# Gradient computation: Stochastic Gradient Descent

Tensorflow  
Federated Learning



# Weight Update

$$\text{update : } \mathbf{w} = \mathbf{w} - \lambda \mathbf{g}$$


Learning rate setting

Update mode





# Weight Update: learning rate




Set the learning rate carefully

# Weight Update: decay learning rate

- ([tensorflow/tensorflow/python/training/learning\\_rate\\_decay.py](https://www.tensorflow.org/tensorflow/python/training/learning_rate_decay.py))

- `exponential_decay`


`exponential_decay(learning_rate, global_step, decay_steps, decay_rate, staircase=False, name=None)`



# Weight Update: decay learning rate


- `piecewise_constant`

`piecewise_constant(x, boundaries, values, name=None)`




# Weight Update: decay learning rate

- polynomial\_decay
- polynomial\_decay(learning\_rate, global\_step, decay\_steps, end\_learning\_rate=0.0001, power=1.0, cycle=False, name=None)



# Weight Update: decay learning rate


- Circle\_decay



Skip local  
optimal

# Weight Update: decay learning rate

- cosine\_decay



# Weight Update : mode

- Gradient descent :

```
for i in range(nb_epochs):  
    params_grad = evaluate_gradient(loss_function, data, params)  
    params = params - learning_rate * params_grad
```

- SGD

```
for i in range(nb_epochs):  
    np.random.shuffle(data)  
    for example in data:  
        params_grad = evaluate_gradient(loss_function, example, params)  
        params = params - learning_rate * params_grad
```



# Weight Update : mode

- Mini-batch Gradient descent

```
for i in range(nb_epochs):  
    np.random.shuffle(data)  
    for batch in get_batches(data, batch_size=50):  
        params_grad = evaluate_gradient(loss_function, batch, params)  
        params = params - learning_rate * params_grad
```

- reduces the variance of the parameter updates, which can lead to more stable convergence
- can make use of highly optimized matrix optimizations common to state-of-the-art deep learning libraries

# Weight Update : mode

- Challenge
  - Choosing a proper learning rate can be difficult
  - Learning rate schedules try to adjust the learning rate during training is defined in advance and are thus unable to adapt to a dataset's characteristics
  - the same learning rate applies to all parameter updates.
  - difficulty arises in fact not from local minima but from saddle points

# Gradient descent optimization algorithms

- Momentum
  - SGD has trouble navigating ravines, i.e. areas where the surface curves much more steeply in one dimension than in another
  - In these scenarios, SGD oscillates across the slopes of the ravine while only making hesitant progress along the bottom towards the local optimum.

# Gradient descent optimization algorithms

- Momentum



Image 2: SGD without momentum




Image 3: SGD with momentum

**Momentum is a method that helps accelerate SGD and dampens oscillations**

# Gradient descent optimization algorithms

- Momentum

$$v_t = \gamma v_{t-1} + \eta \nabla_{\theta} J(\theta)$$
$$\theta = \theta - v_t$$

Momentum

Fraction = 0.9

Last time update vector

*If we push a ball down a hill. The ball accumulates momentum as it rolls downhill, becoming faster and faster on the way*

Why effective

- The momentum term increases for dimensions whose gradients point in the same directions
- reduces updates for dimensions whose gradients change directions.
- As a result, we gain faster convergence and reduced oscillation.

# Gradient descent optimization algorithms

- Nesterov accelerated gradient (NAG)
  - a ball that rolls down a hill blindly, is highly unsatisfactory
  - We need **a smarter** ball, a ball that has a notion of where it is going to

momentum  $\longrightarrow v_t = \gamma v_{t-1} + \eta \nabla_{\theta} J(\theta)$

Gradient w.r.t the current parameters

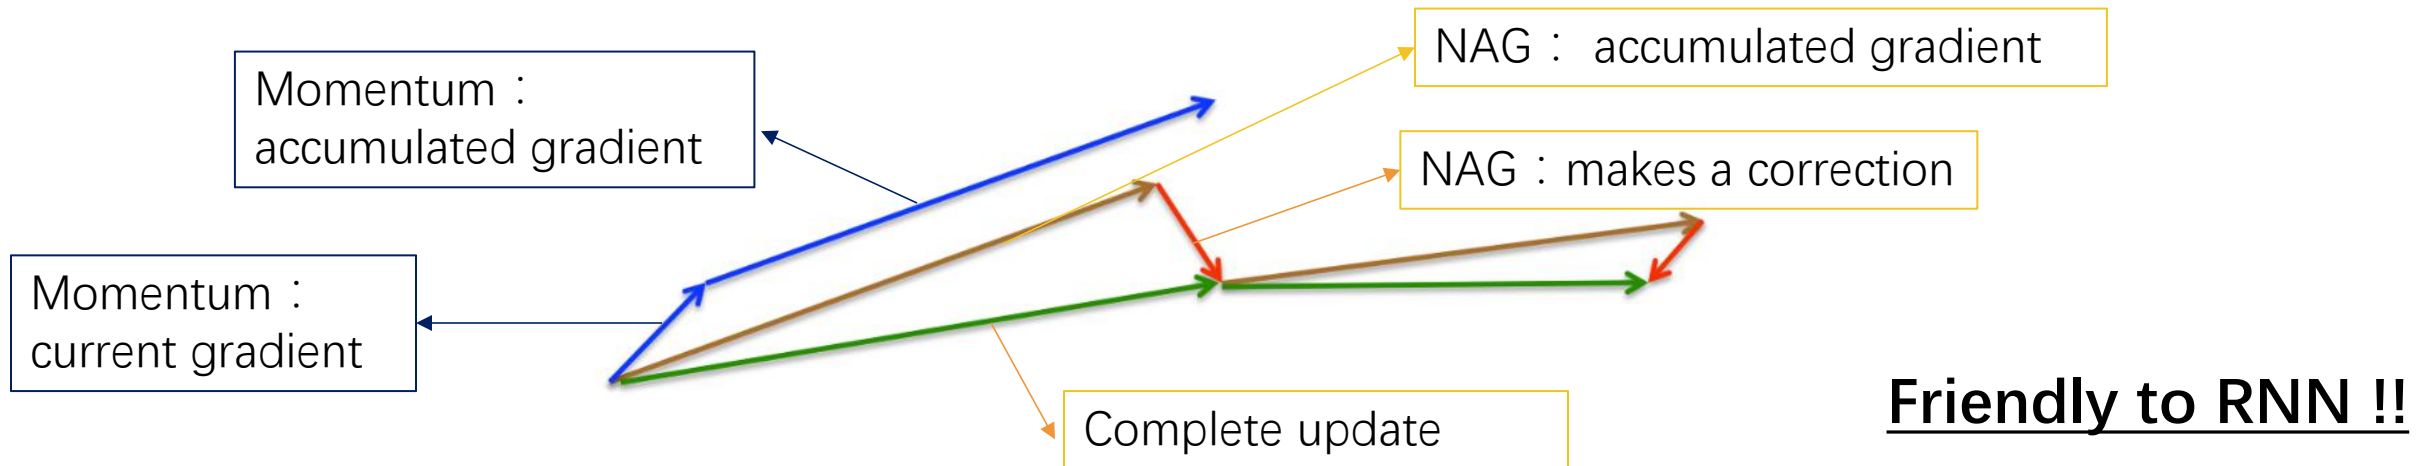
NAG  $\longrightarrow$   
$$v_t = \gamma v_{t-1} + \eta \nabla_{\theta} J(\theta - \gamma v_{t-1})$$
$$\theta = \theta - v_t$$

Gradient w.r.t approximate future position of our parameters

*NAG can now effectively look ahead by calculating the gradient not w.r.t. to our current parameters but w.r.t. the approximate future position of our parameters*

# Gradient descent optimization algorithms

- Nesterov accelerated gradient (NAG)
  - a ball that rolls down a hill blindly, is highly unsatisfactory
  - We need **a smarter** ball, a ball that has a notion of where it is going to



*NAG can now effectively look ahead by calculating the gradient not w.r.t. to our current parameters but w.r.t. the approximate future position of our parameters*



# Gradient descent optimization algorithms

- Adagrad uses a different learning rate for every parameter at every time step
- Adagrad adapts the learning rate to the parameters,
  - **smaller updates (i.e. low learning rates)** for parameters associated with frequently occurring features,
  - **larger updates (i.e. high learning rates)** for parameters associated with infrequent features

# Gradient descent optimization algorithms

## SGD

$$g_{t,i} = \nabla_{\theta} J(\theta_{t,i}).$$

$$\theta_{t+1,i} = \theta_{t,i} - \eta \cdot g_{t,i}.$$

## Adagrad

$$g_{t,i} = \nabla_{\theta} J(\theta_{t,i}).$$

$$\theta_{t+1,i} = \theta_{t,i} - \frac{\eta}{\sqrt{G_{t,ii} + \epsilon}} \cdot g_{t,i}.$$

diagonal matrix where each diagonal element is the sum of the squares of the gradients w.r.t. up to time step  $t$ .

# Gradient descent optimization algorithms

- adagrad
  - Benefits
    - it eliminates the need to manually tune the learning rate
  - Weakness:
    - the learning rate to shrink and eventually become infinitesimally small

# Gradient descent optimization algorithms

- Adaptive Moment Estimation (Adam)
  - is another method that computes adaptive learning rates for each parameter

$$m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t$$
$$v_t = \beta_2 v_{t-1} + (1 - \beta_2) g_t^2$$

$m$  and  $v$  are estimates of the first moment (the mean) and the second moment (the uncentered variance) of the gradients



$$\hat{m}_t = \frac{m_t}{1 - \beta_1^t}$$
$$\hat{v}_t = \frac{v_t}{1 - \beta_2^t}$$

counteract these biases



$$\theta_{t+1} = \theta_t - \frac{\eta}{\sqrt{\hat{v}_t} + \epsilon} \hat{m}_t.$$

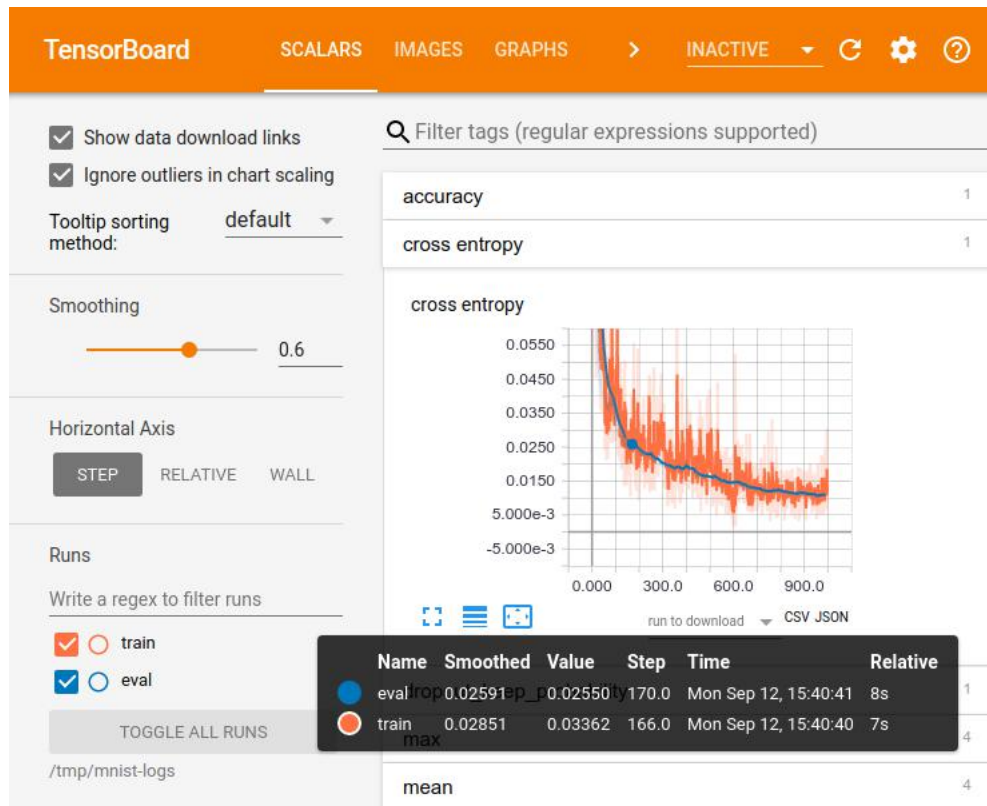
Final update ! ! !

# Code example

Image classification in colab, and hyperparameter tuning using tensorboard

# Example

- Tensorboard
  - Tensorboard visualization tools



- Can also be used in Colab & Google Cloud
- New features: hyperparameter tuning in TF2.0

2019  
怪兽  
学堂



THANKS