

2019
怪兽
学堂

GAN



虾米

2019-04

Outline

- **Part 1: Introduction to GANs**
- **Part 2: Some challenges with GANs**
- **Part 3: Applications of GANs**

Part 1

- Motivation for Generative Models
- From Adversarial Training to GANs
- GAN's Architecture
- GAN's objective
- DCGANs

GANs

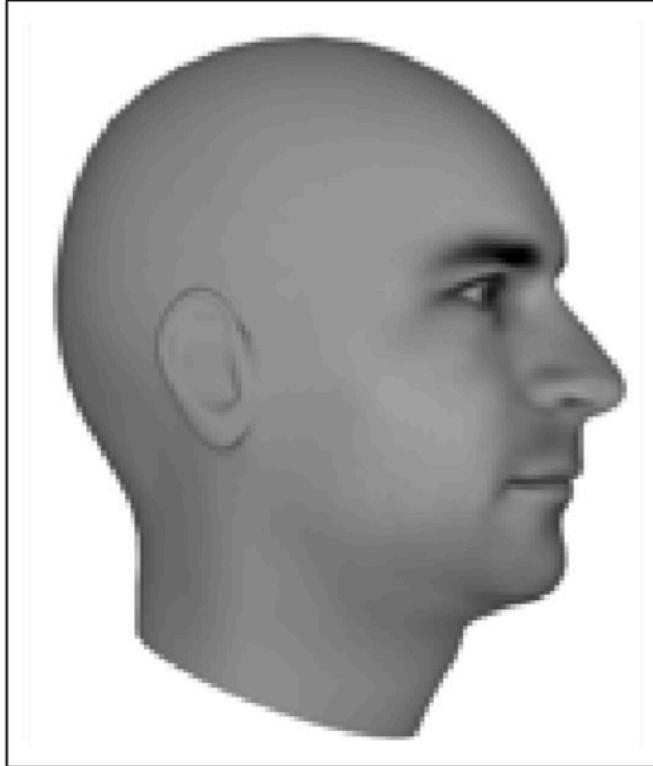
- **Generative**
 - Learn a generative model
- **Adversarial**
 - Trained in an adversarial setting
- **Networks**
 - Use Deep Neural Networks

Why Generative Models?

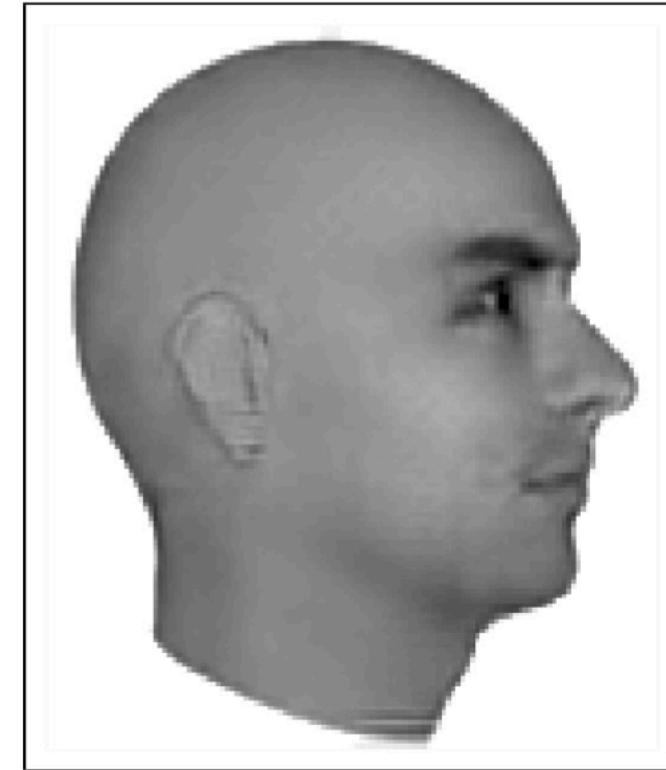
- **We've only seen discriminative models so far**
 - Given an image X , predict a label Y
 - Estimates $P(Y|X)$
- **Discriminative models have several key limitations**
 - Can't model $P(X)$, i.e. the probability of seeing a certain image
 - Thus, can't sample from $P(X)$, i.e. **can't generate new images**
- **Generative models (in general) cope with all of above**
 - Can model $P(X)$
 - Can generate new images

Magic of GANs...

Ground Truth

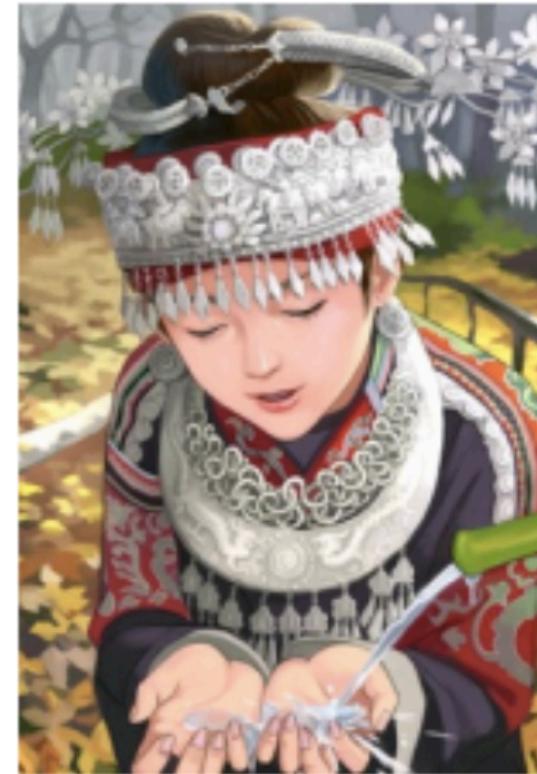


Adversarial



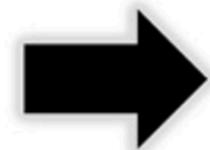
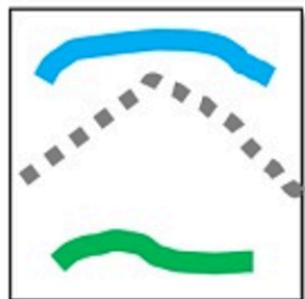
Magic of GANs...

Which one is Computer generated?



Magic of GANs...

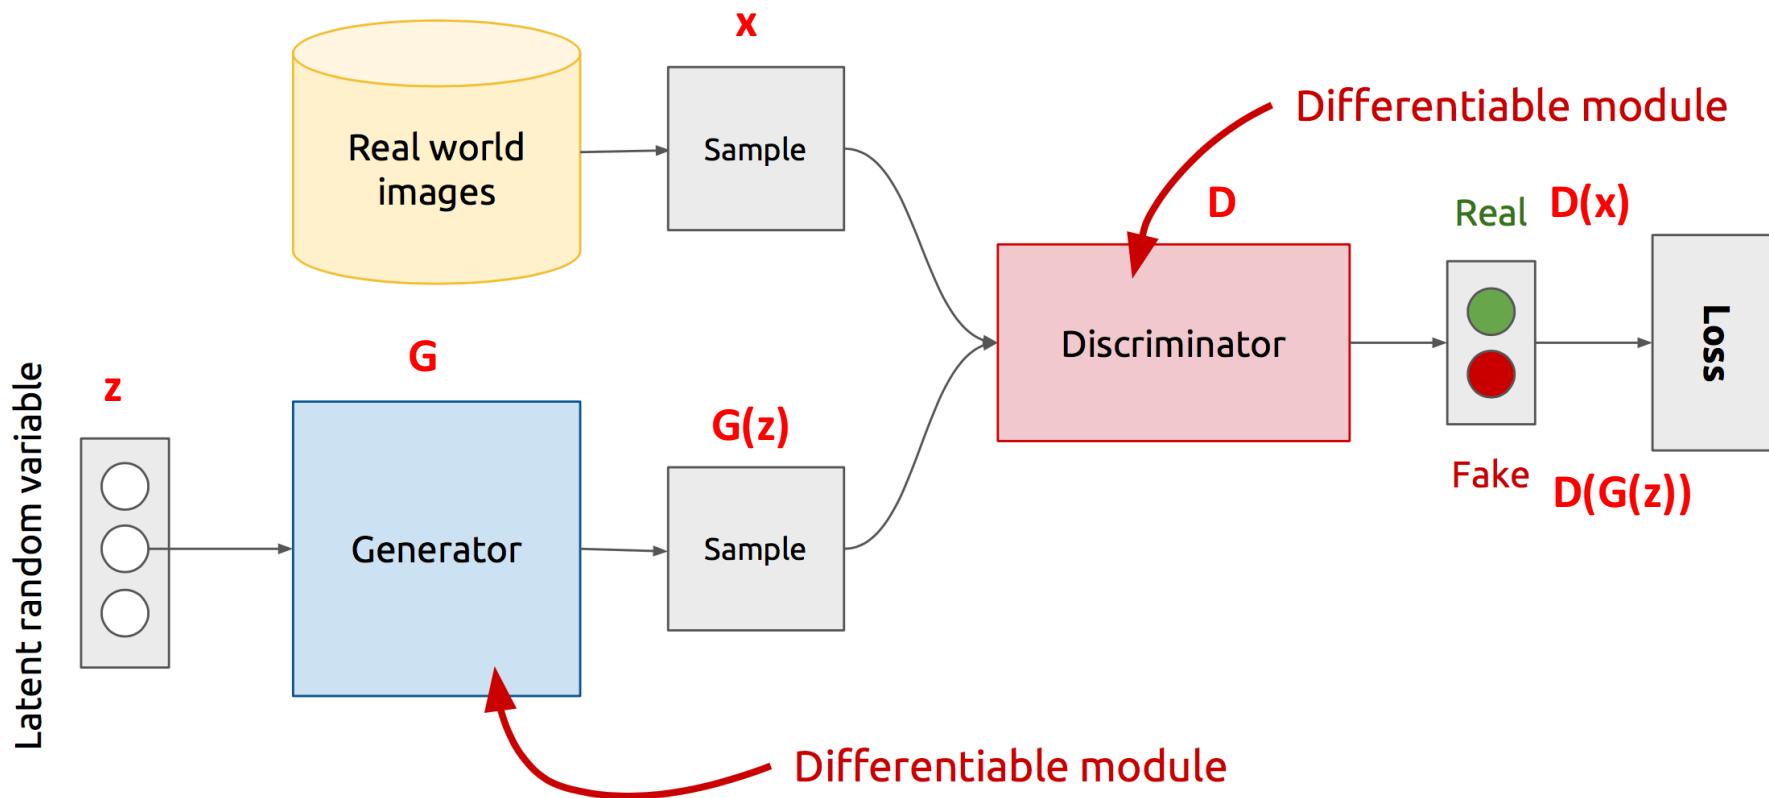
User edits



Generated images

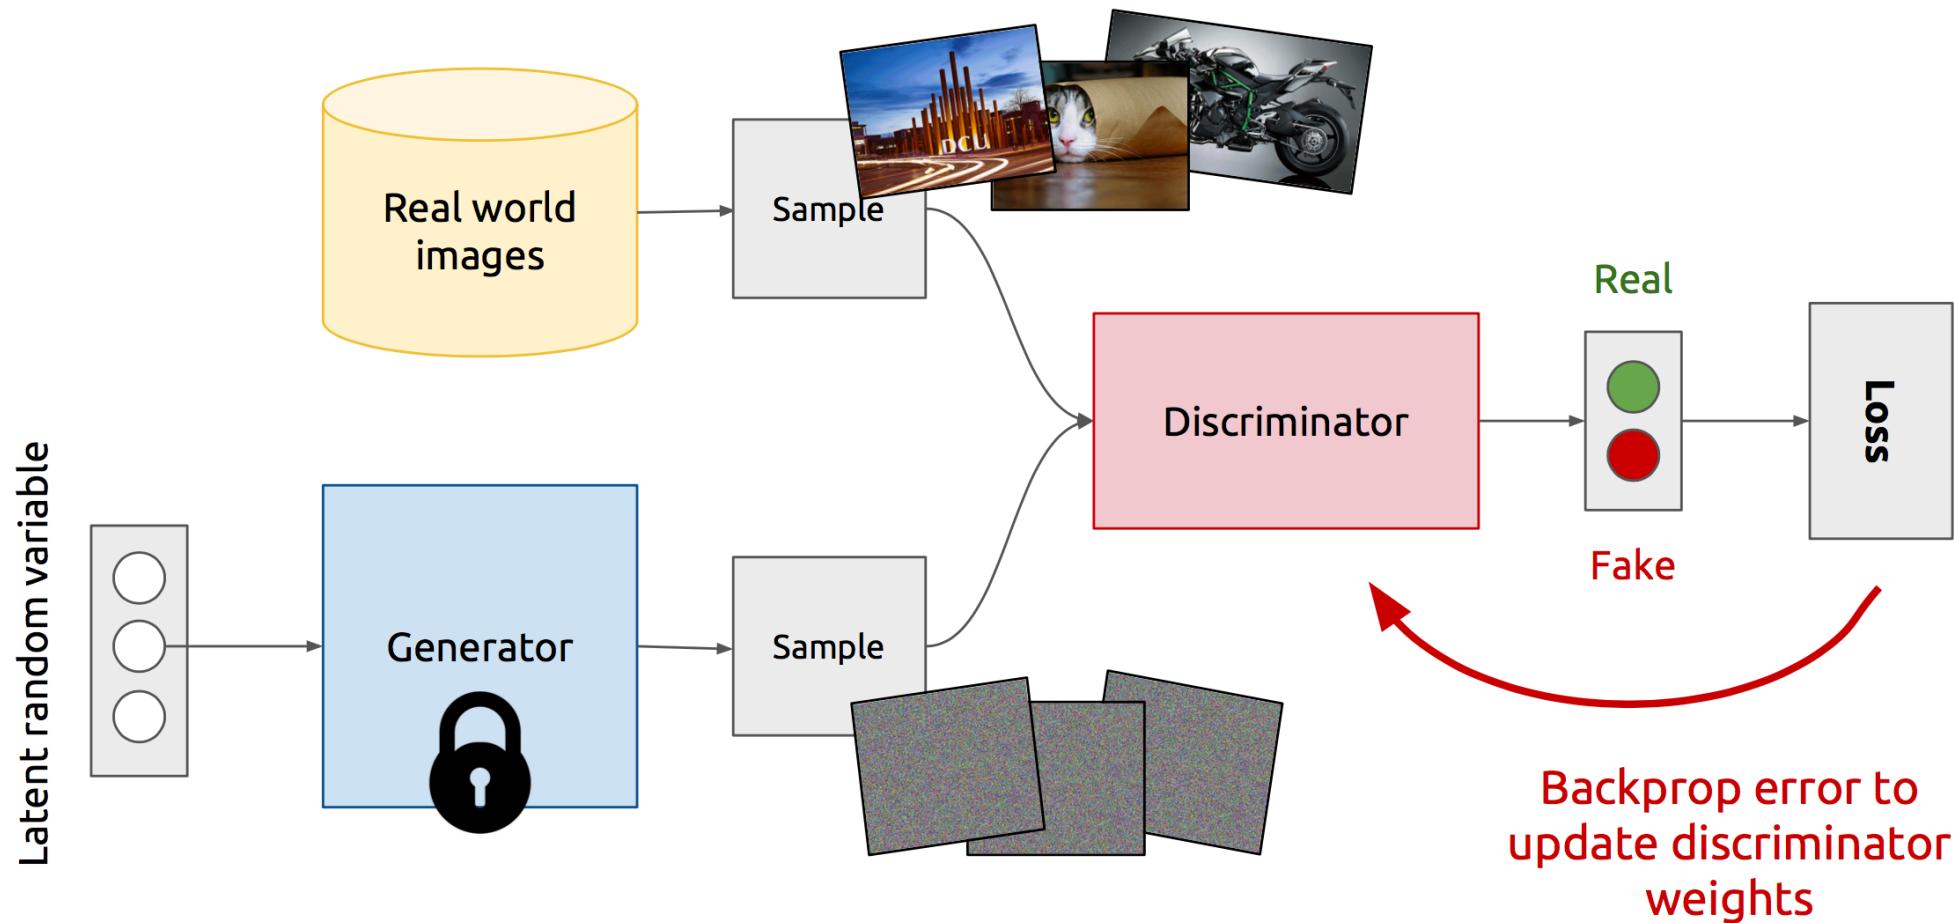


GAN's Architecture

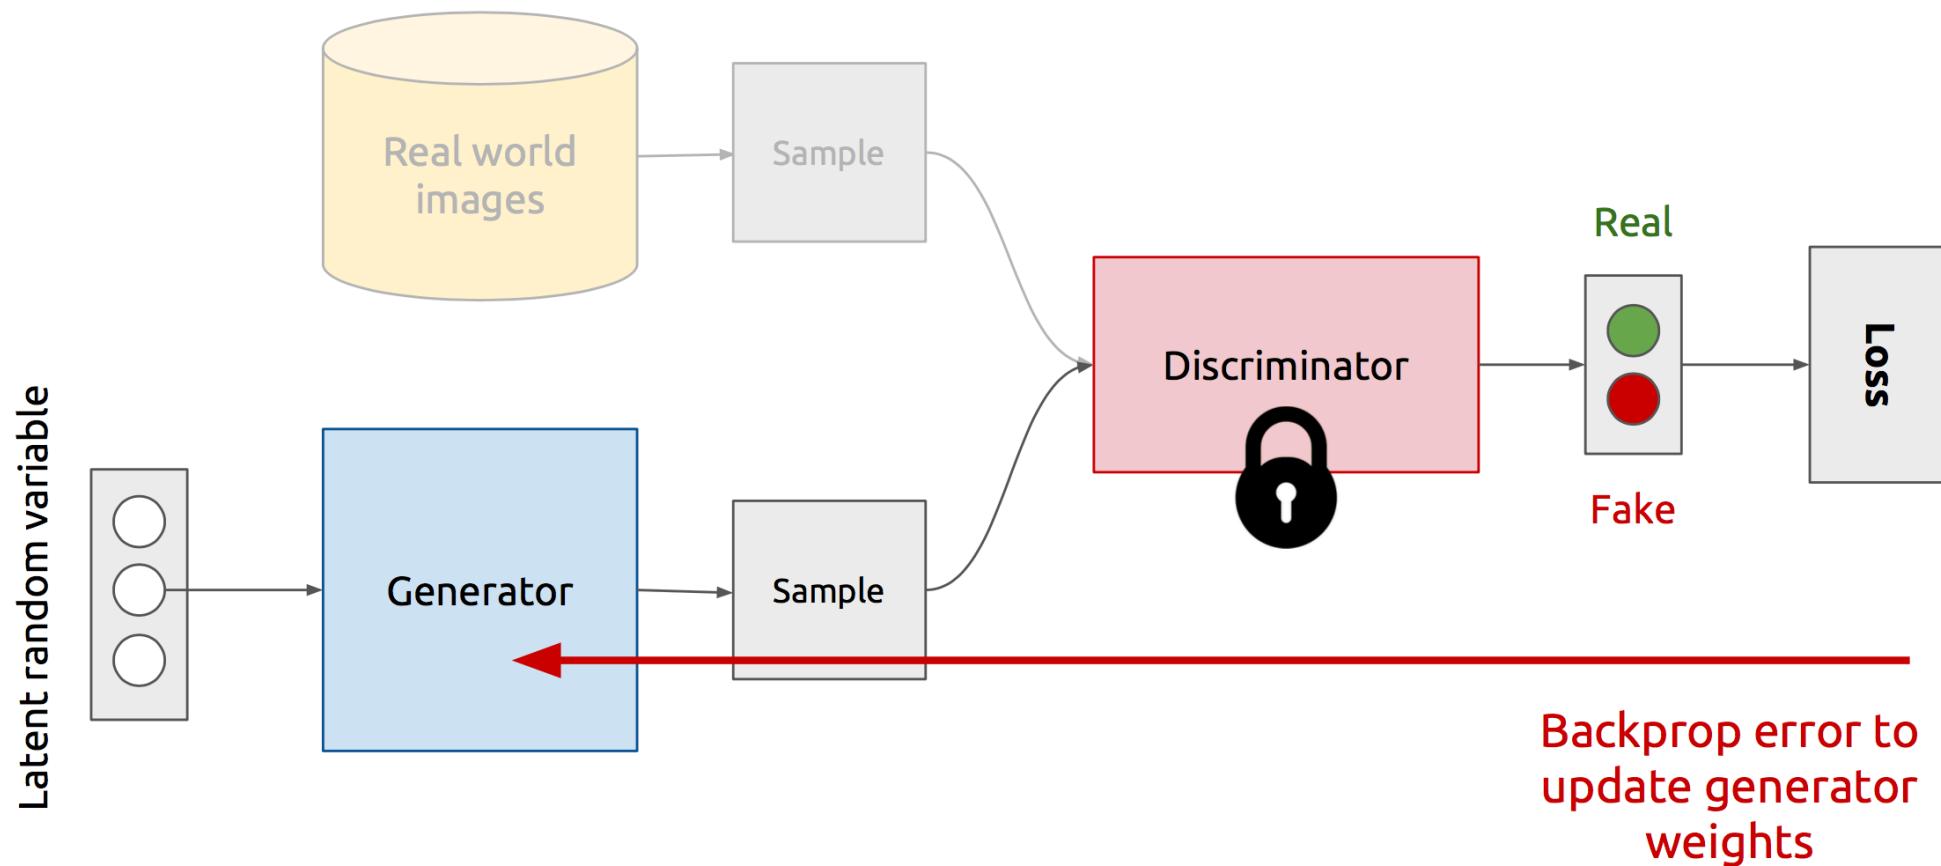


- Z is some random noise (Gaussian/Uniform).
- Z can be thought as the latent representation of the image.

Training Discriminator



Training Generator



GAN's formulation

$$\min_G \max_D V(D, G)$$

- It is formulated as a **minimax game**, where:
 - The Discriminator is trying to maximize its reward $V(D, G)$
 - The Generator is trying to minimize Discriminator's reward (or maximize its loss)

$$V(D, G) = \boxed{\mathbb{E}_{x \sim p(x)} [\log D(x)]} + \boxed{\mathbb{E}_{z \sim q(z)} [\log(1 - D(G(z)))]}$$

- The Nash equilibrium of this particular game is achieved at:
 - $P_{data}(x) = P_{gen}(x) \quad \forall x$
 - $D(x) = \frac{1}{2} \quad \forall x$

**Discriminator
updates**

Algorithm 1 Minibatch stochastic gradient descent training of generative adversarial nets. The number of steps to apply to the discriminator, k , is a hyperparameter. We used $k = 1$, the least expensive option, in our experiments.

for number of training iterations **do**

for k steps **do**

- Sample minibatch of m noise samples $\{z^{(1)}, \dots, z^{(m)}\}$ from noise prior $p_g(z)$.
- Sample minibatch of m examples $\{x^{(1)}, \dots, x^{(m)}\}$ from data generating distribution $p_{\text{data}}(x)$.
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[\log D(x^{(i)}) + \log (1 - D(G(z^{(i)}))) \right].$$

end for

- Sample minibatch of m noise samples $\{z^{(1)}, \dots, z^{(m)}\}$ from noise prior $p_g(z)$.
- Update the generator by descending its stochastic gradient:

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log (1 - D(G(z^{(i)}))).$$

end for

The gradient-based updates can use any standard gradient-based learning rule. We used momentum in our experiments.

**Generator
updates**

Vanishing gradient strikes back again...

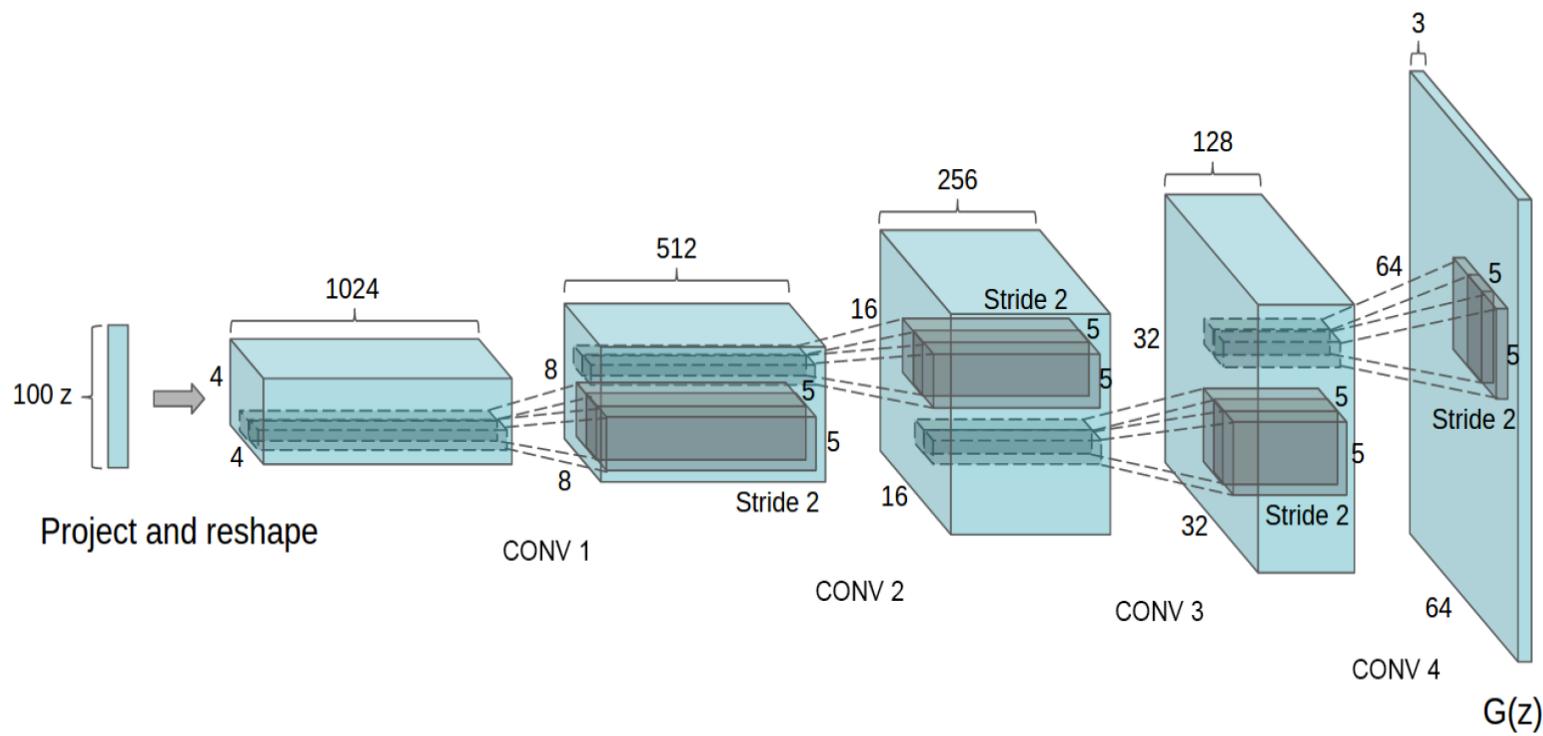
$$V(D, G) = \min_G \max_D V(D, G)$$
$$V(D, G) = \mathbb{E}_{x \sim p(x)} [\log D(x)] + \boxed{\mathbb{E}_{z \sim q(z)} [\log(1 - D(G(z)))]}$$

$$\nabla_{\theta_G} V(D, G) = \nabla_{\theta_G} \mathbb{E}_{z \sim q(z)} [\log(1 - D(G(z)))]$$

- $\nabla_a \log(1 - \sigma(a)) = \frac{-\nabla_a \sigma(a)}{1 - \sigma(a)} = \frac{-\sigma(a)(1 - \sigma(a))}{1 - \sigma(a)} = -\sigma(a) = -D(G(z))$
- Gradient goes to 0 if D is confident, i.e. $D(G(z)) \rightarrow 0$
- Minimize $-\mathbb{E}_{z \sim q(z)} [\log D(G(z))]$ for **Generator** instead (keep Discriminator as it is)

Deep Convolutional GANs (DCGANs)

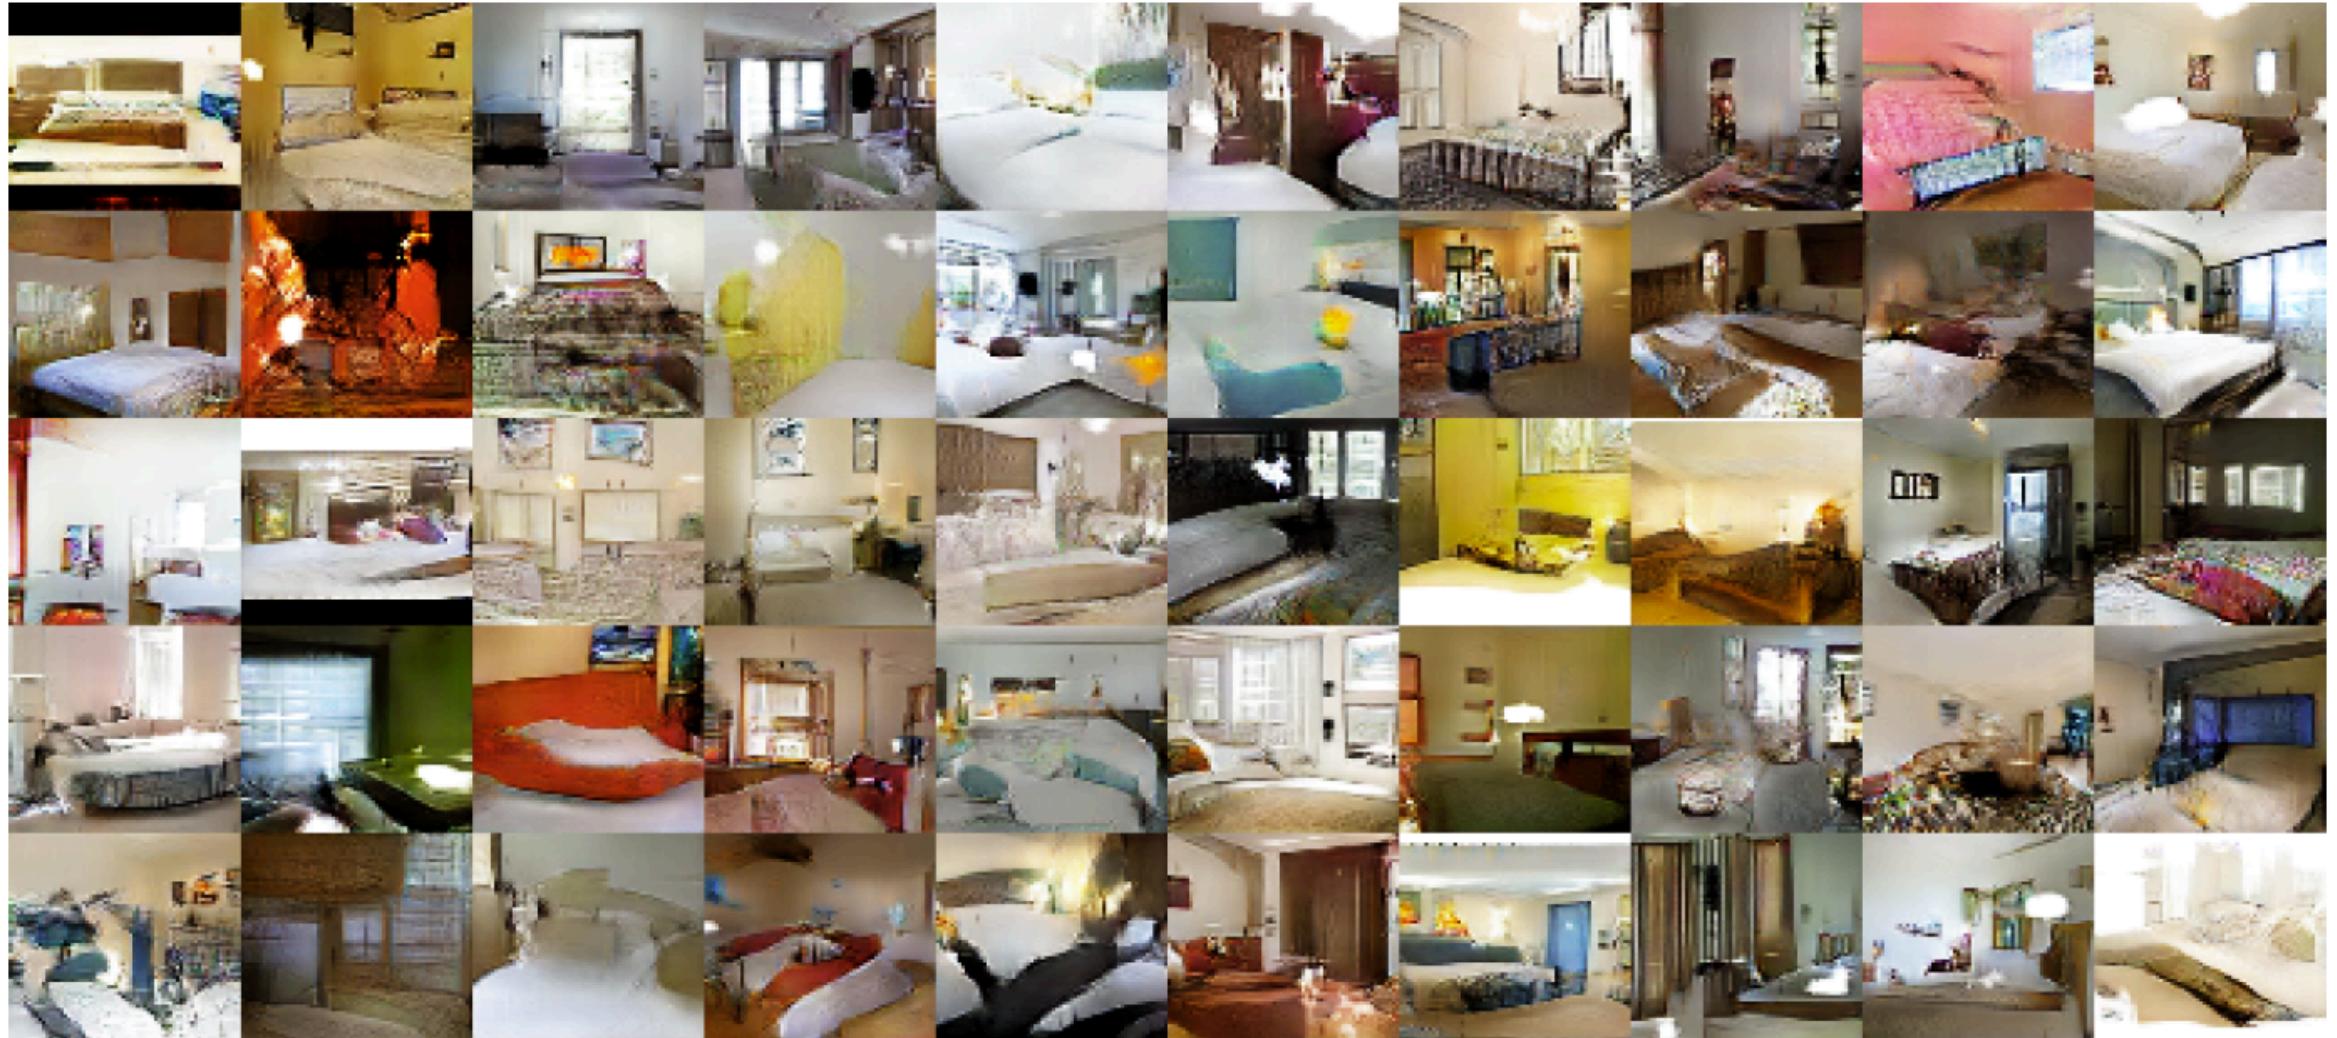
Generator Architecture



Key ideas:

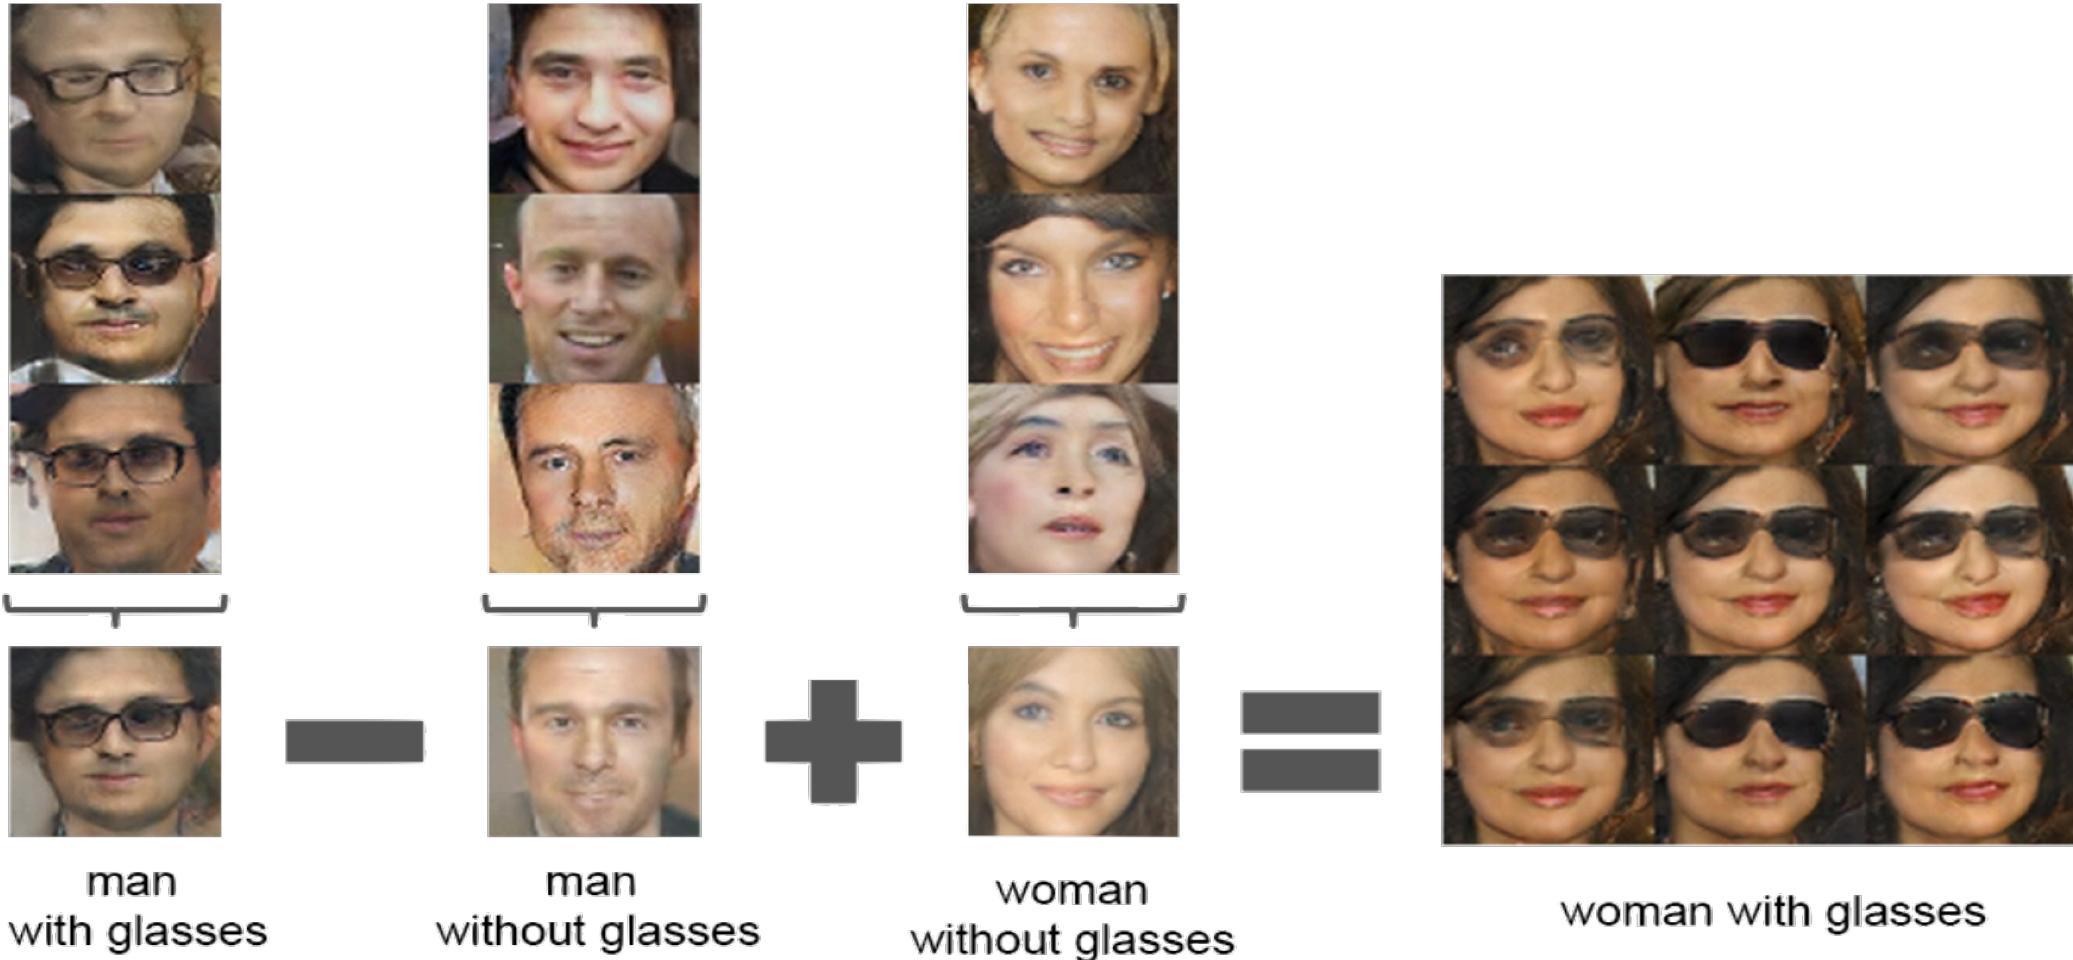
- Replace FC hidden layers with Convolutions
 - **Generator:** Fractional-Strided convolutions
- Use Batch Normalization after each layer
- **Inside Generator**
 - Use ReLU for hidden layers
 - Use Tanh for the output layer

DCGAN: Bedroom images



Radford, Alec, Luke Metz, and Soumith Chintala. "Unsupervised representation learning with deep convolutional generative adversarial networks." arXiv:1511.06434 (2015).

Latent vectors capture interesting patterns...



Part 2

- **Advantages of GANs**
- **Training Challenges**
 - Non-Convergence
 - Mode-Collapse
- **Proposed Solutions**
 - Supervision with Labels
 - Mini-Batch GANs
- **Modification of GAN's losses**
 - Discriminator (**EB-GAN**)
 - Generator (**InfoGAN**)

Advantages of GANs

- **Plenty of existing work on Deep Generative Models**

- Boltzmann Machine
- Deep Belief Nets
- Variational AutoEncoders (VAE)

- **Why GANs?**

- Sampling (or generation) is straightforward.
- Training doesn't involve Maximum Likelihood estimation.
- Robust to Overfitting since Generator never sees the training data.
- Empirically, GANs are good at capturing the modes of the distribution.

Problems with GANs

- **Probability Distribution is Implicit**
 - Not straightforward to compute $P(X)$.
 - Thus **Vanilla GANs** are only good for Sampling/Generation.
- **Training is Hard**
 - Non-Convergence
 - Mode-Collapse

Training Problems

- **Non-Convergence**
- Mode-Collapse

- Deep Learning models (in general) involve a single player

- The player tries to maximize its reward (minimize its loss).
- Use SGD (with Backpropagation) to find the optimal parameters.
- SGD has convergence guarantees (under certain conditions).
- **Problem:** With non-convexity, we might converge to local optima.

$$\min_G L(G)$$

- GANs instead involve two (or more) players

- Discriminator is trying to maximize its reward.
- Generator is trying to minimize Discriminator's reward.

$$\min_G \max_D V(D, G)$$

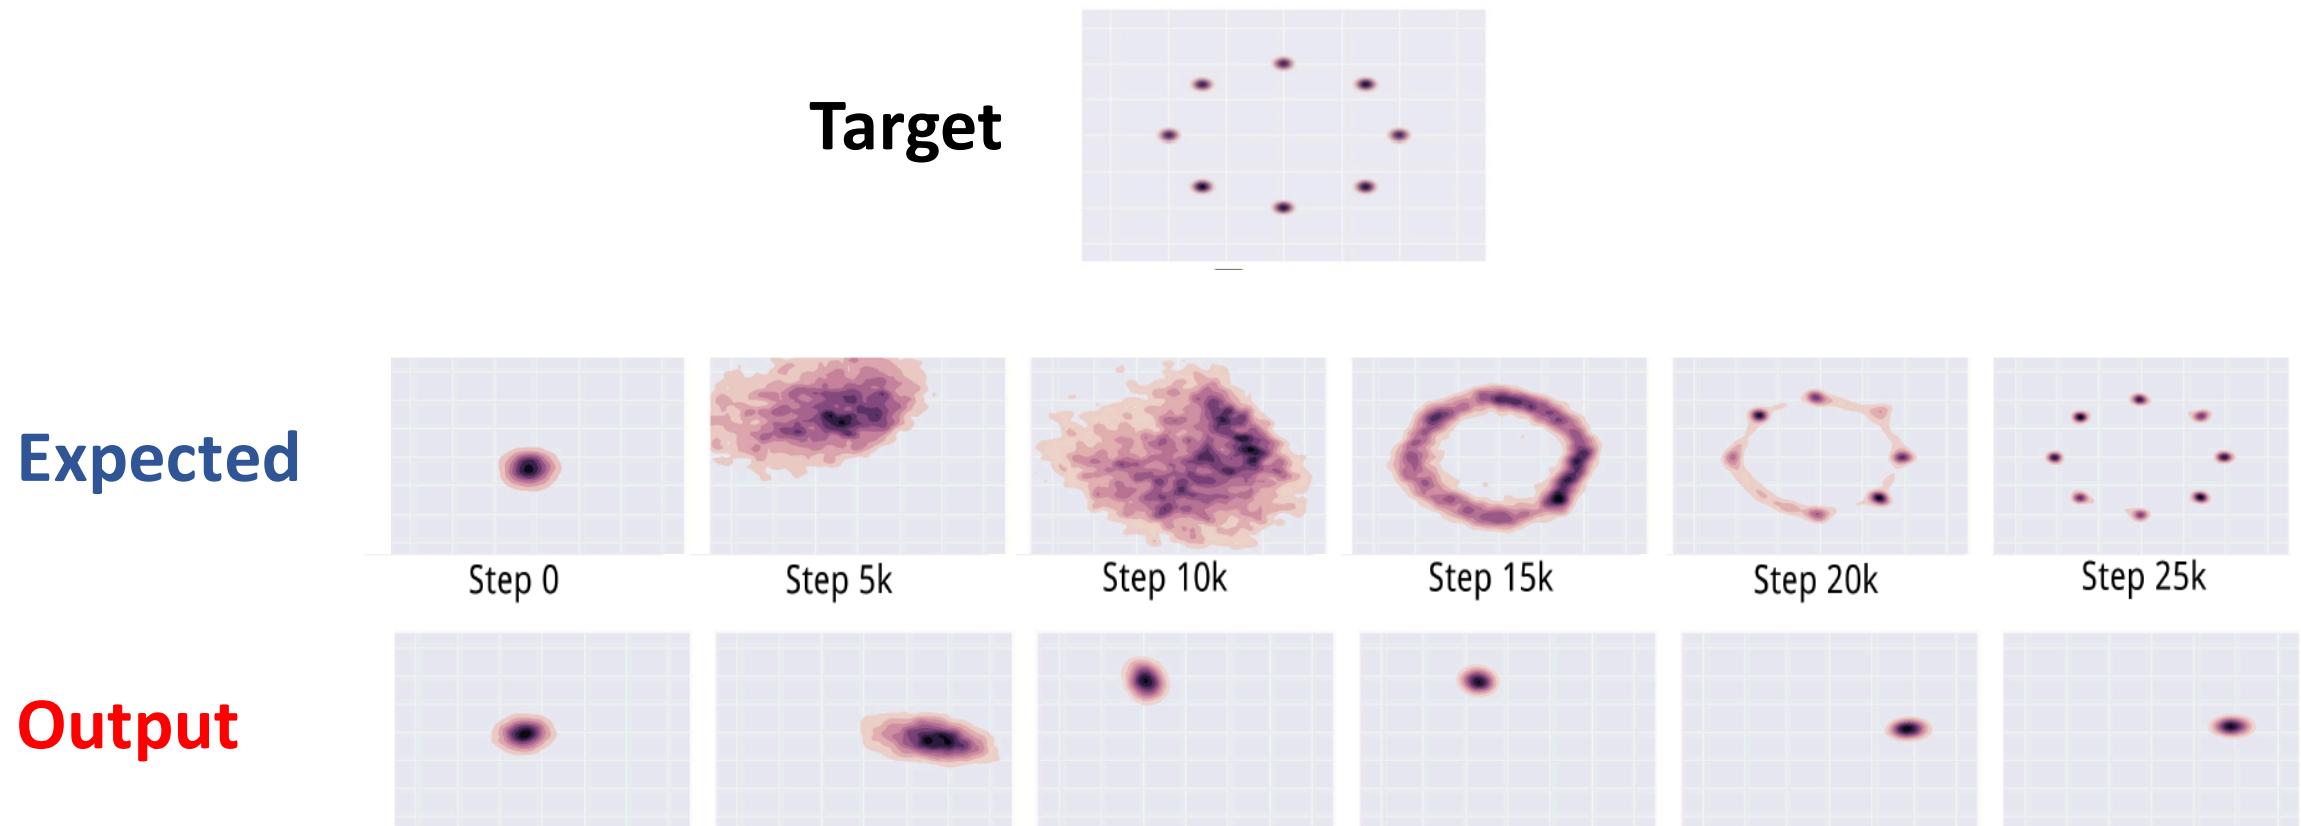
- SGD was not designed to find the Nash equilibrium of a game.
- **Problem:** We might not converge to the Nash equilibrium at all.

Problems with GANs

- Non-Convergence
- **Mode-Collapse**

Mode-Collapse

- Generator fails to output diverse samples



Some real examples



Some Solutions

- Mini-Batch GANs
- Supervision with labels
- Some recent attempts :-
 - Unrolled GANs
 - W-GANs

Basic (Heuristic) Solutions

- **Mini-Batch GANs**
- Supervision with labels

How to reward sample diversity?

- **At Mode Collapse,**
 - Generator produces good samples, but a very few of them.
 - Thus, Discriminator can't tag them as fake.
- **To address this problem,**
 - Let the Discriminator know about this edge-case.
- **More formally,**
 - Let the Discriminator look at the entire batch instead of single examples
 - If there is lack of diversity, it will mark the examples as fake
- **Thus,**
 - Generator will be forced to produce diverse samples.

Mini-Batch GANs

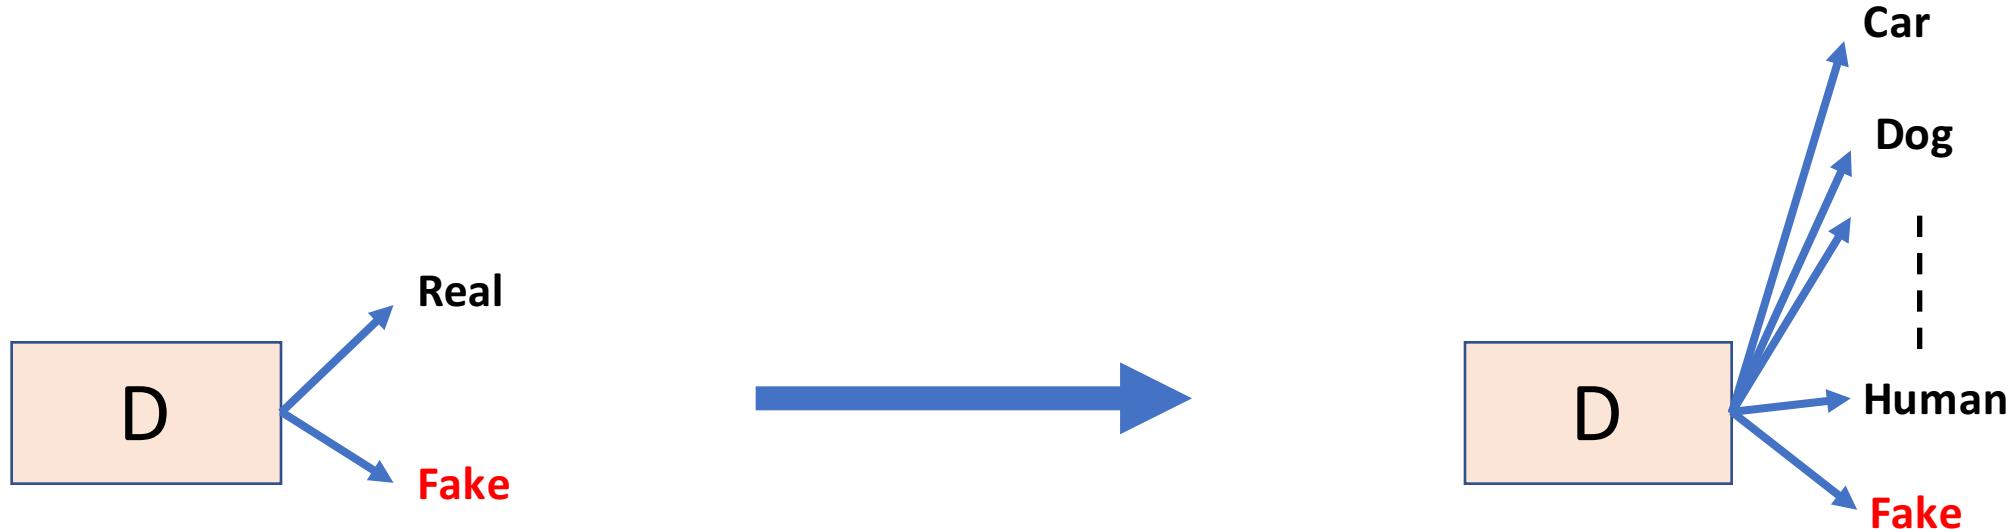
- Extract features that capture diversity in the mini-batch
 - For e.g. L2 norm of the difference between all pairs from the batch
- Feed those features to the discriminator along with the image
- Feature values will differ b/w diverse and non-diverse batches
 - Thus, Discriminator will rely on those features for classification
- This in turn,
 - Will force the Generator to match those feature values with the real data
 - Will generate diverse batches

Basic (Heuristic) Solutions

- Mini-Batch GANs
- **Supervision with labels**

Supervision with Labels

- Label information of the real data might help



- Empirically generates much better samples

Part 3

- **Conditional GANs**
- **Applications**
 - Image-to-Image Translation
 - Text-to-Image Synthesis
 - Face Aging
- **Advanced GAN Extensions**
 - Coupled GAN
 - LAPGAN – Laplacian Pyramid of Adversarial Networks
 - Adversarially Learned Inference
- **Summary**

Conditional GANs

MNIST digits generated conditioned on their class label.

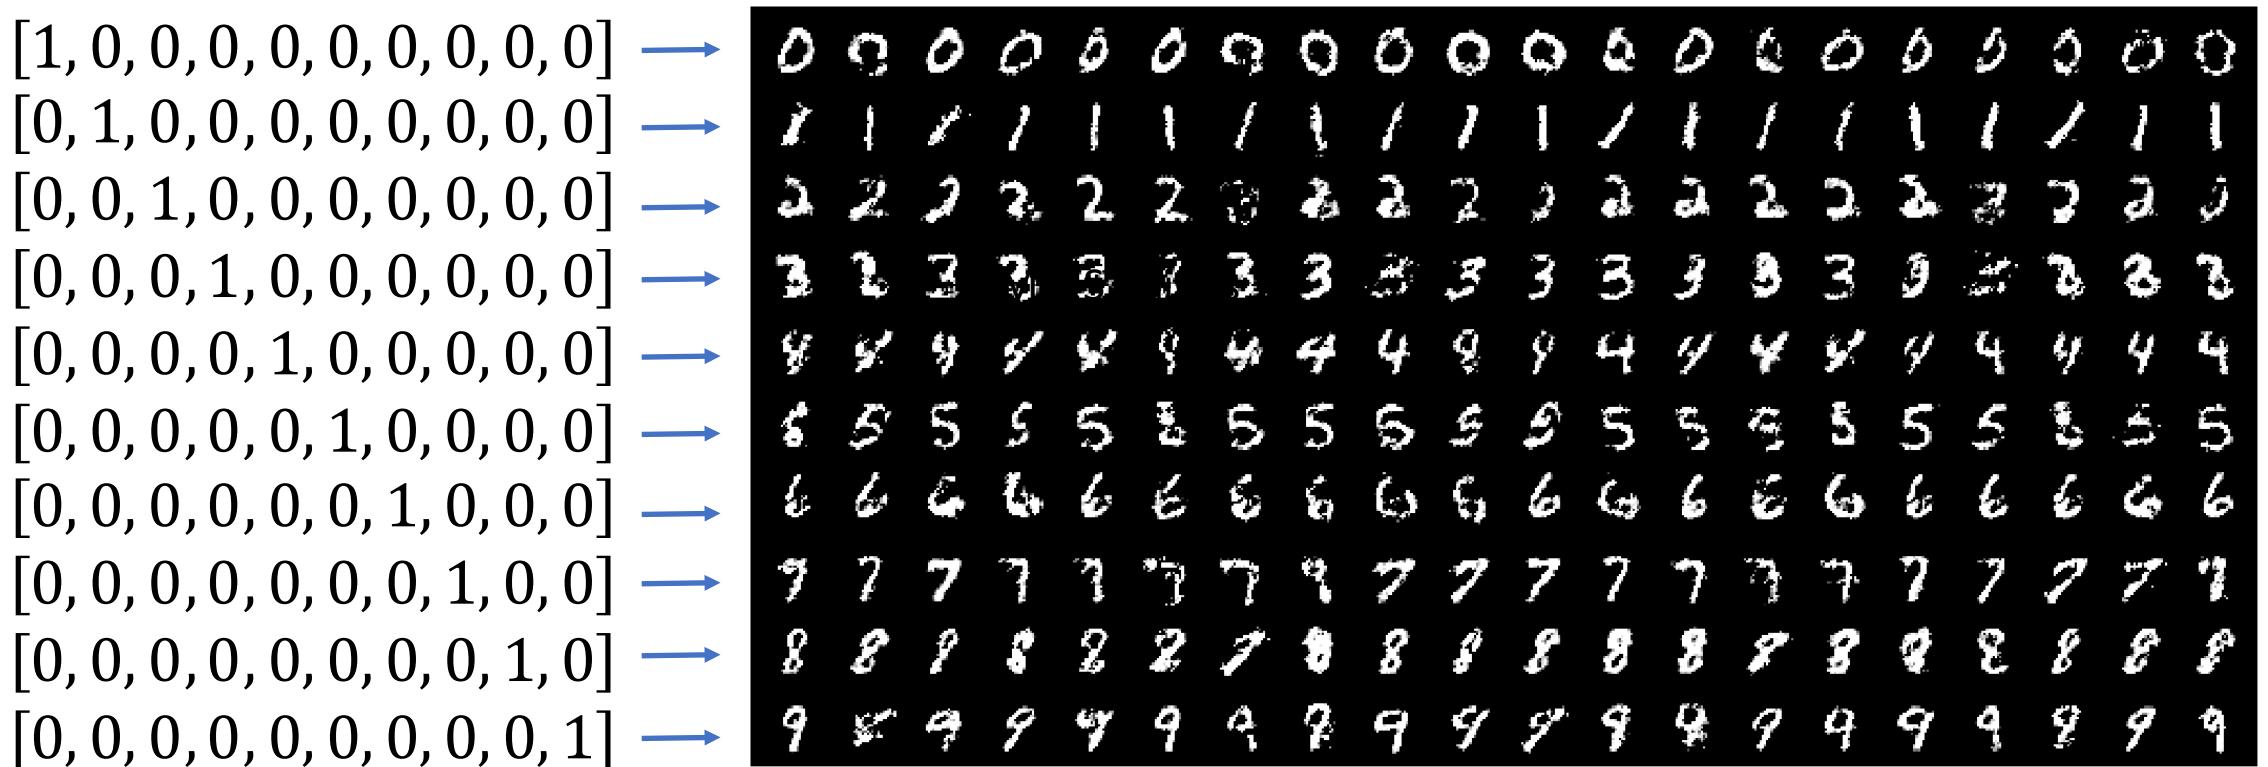
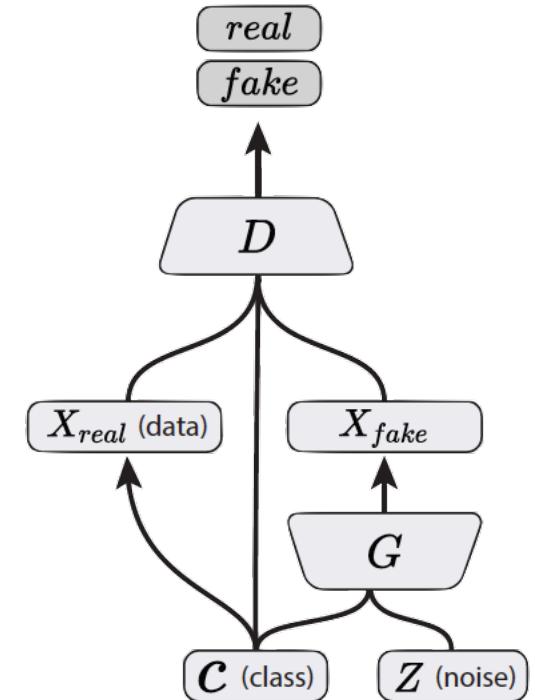


Figure 2 in the original paper.

Conditional GANs

- Simple modification to the original GAN framework that conditions the model on *additional information* for better multi-modal learning.
- Lends to many practical applications of GANs when we have explicit *supervision* available.



Conditional GAN
(Mirza & Osindero, 2014)

Image Credit: Figure 2 in Odena, A., Olah, C. and Shlens, J., 2016. Conditional image synthesis with auxiliary classifier GANs. *arXiv preprint arXiv:1610.09585*.

Part 3

- **Conditional GANs**
- **Applications**
 - Image-to-Image Translation
 - Text-to-Image Synthesis
 - Face Aging
- **Advanced GAN Extensions**
 - Coupled GAN
 - LAPGAN – Laplacian Pyramid of Adversarial Networks
 - Adversarially Learned Inference
- **Summary**

Image-to-Image Translation

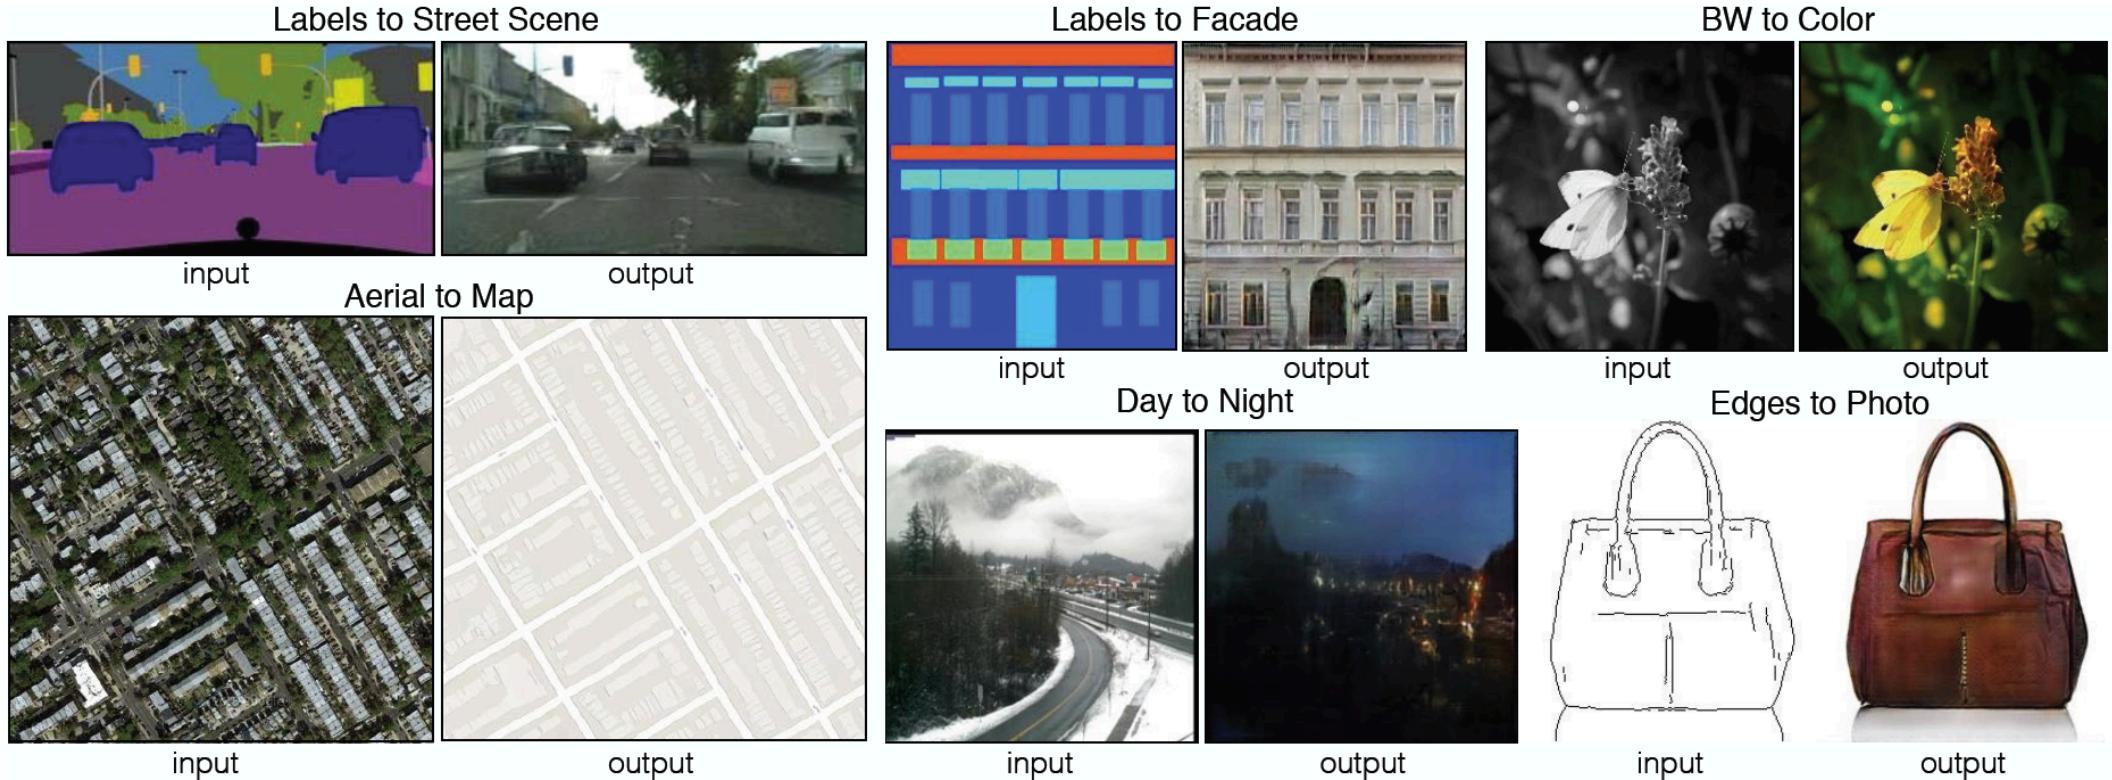


Figure 1 in the original paper.

[Link to an interactive demo of this paper](#)

Image-to-Image Translation

- Architecture: *DCGAN-based architecture*
- Training is conditioned on the images from the source domain.
- Conditional GANs provide an effective way to handle many complex domains without worrying about designing *structured loss* functions explicitly.

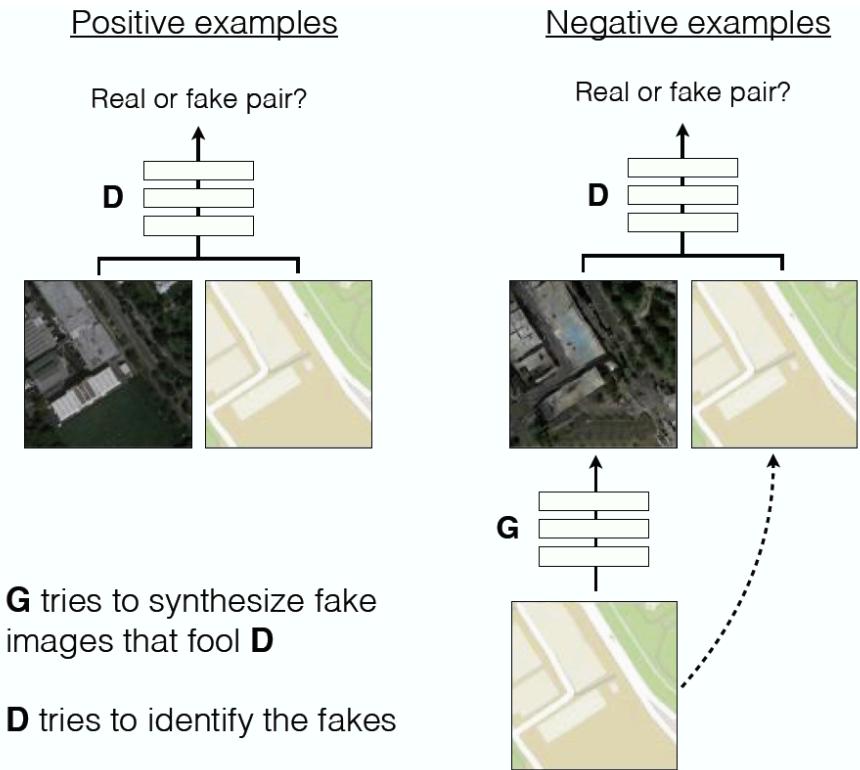


Figure 2 in the original paper.

Conditional GANs

Conditional Model Collapse

- Scenario observed when the Conditional GAN starts *ignoring* either the code (c) or the noise variables (z).
- This limits the diversity of images generated.



Credit?

Part 3

- **Conditional GANs**
- **Applications**
 - Image-to-Image Translation
 - Text-to-Image Synthesis
 - Face Aging
- **Advanced GAN Extensions**
 - Coupled GAN
 - LAPGAN – Laplacian Pyramid of Adversarial Networks
 - Adversarially Learned Inference
- **Summary**

Summary

- GANs are generative models that are implemented using two stochastic neural network modules: **Generator** and **Discriminator**.
- **Generator** tries to generate samples from random noise as input
- **Discriminator** tries to distinguish the samples from Generator and samples from the real data distribution.
- Both networks are trained adversarially (in tandem) to fool the other component. In this process, both models become better at their respective tasks.

补充材料

- https://mp.weixin.qq.com/s?__biz=MzI5ODMzMjk0NA==&mid=100000106&idx=1&sn=6d2440716ad5cffa9535c785bbf71372&chksm=6ca627155bd1ae0302c02e6a34f3437b992e68221c30e83d1fc8817952215195b4bd815adca6&mpshare=1&scene=1&srcid=#rd

2019
怪兽
学堂

THANKS

